

Lakshmi Sahasra Jangoan

Atlanta, GA | +1 (405)-659-0961 | jangoanlakshmisahasra@gmail.com | LinkedIn | GitHub | Portfolio

PROFESSIONAL SUMMARY

Data Professional with 2.5+ years of experience turning large, messy datasets into reliable ML models, BI dashboards, ETL workflows, and stakeholder-ready insights. Recently completed an M.S. in Data Science and Analytics with 4.0 GPA, after building research and business analytics solutions across climate forecasting, healthcare, telecom, mobility, AI assistants, and operational reporting. Strong in SQL, Python, R, Power BI, Tableau, statistics, machine learning, time series, data engineering, and GenAI/RAG, with proven impact including 14% forecast accuracy lift, 18% RMSE reduction, 25% reporting automation savings, and 15% reporting accuracy improvement.

TECHNICAL SKILLS

Languages & Querying: Python, SQL, R, Bash, JavaScript, MATLAB, Java, CTEs, Joins, Window Functions, Stored Procedures, Aggregations, Subqueries

Data Analytics & BI: Excel, Power BI, Tableau, DAX, Power Query, KPI Reporting, Dashboard Development, Data Visualization, Data Storytelling, Ad-hoc Analysis, EDA

Machine Learning & Statistics: Regression, Classification, Clustering, XGBoost, Random Forest, Logistic Regression, LSTM, ARIMA, Forecasting, Hypothesis Testing, Probability, Cross-Validation, Model Evaluation

AI & NLP: PyTorch, TensorFlow/Keras, Hugging Face Transformers, NLP, LLMs, RAG, Embeddings, Prompt Engineering, SHAP Explainability

Data Engineering & Tools: PostgreSQL, Azure SQL, Databricks, Delta Lake, dbt, Airflow, FastAPI, Docker, REST APIs, Pytest, Git, GitHub, Data Cleaning, Data Validation, ETL/ELT

PROFESSIONAL EXPERIENCE

Data Scientist - Research | *National Weather Center, University of Oklahoma* | *Norman, OK* Oct 2024 – Jun 2026

- Faced with noisy, nonstationary climate data, built automated Python ML pipelines to clean, transform, and model 40+ years of ERA5 reanalysis data across 156 atmospheric channels.
- Combined ResNet18 feature extraction with LSTM, autoregressive, Random Forest, Logistic Regression, and XGBoost models, improving regional precipitation forecast accuracy by 14%.
- Engineered temporal lag/lead and spatiotemporal features across 600+ HPC experiments, reducing RMSE by 18% and MAE by 13% through model tuning and error analysis.
- Validated reliability with calibration, anomaly detection, and distribution-aware diagnostics, then translated results into visual summaries for research stakeholders.

Data Analyst | *ValueLabs* | *Hyderabad, India* Nov 2023 – Jul 2024

- To improve operational visibility, analyzed 120,000+ delivery, resource utilization, and billing records using SQL, Python, Pandas, Excel, and Databricks to uncover KPI trends and delays.
- Built Power BI dashboards tracking 15+ KPIs with DAX, Power Query, slicers, and drill-down views, giving leadership clearer visibility into project performance and resource utilization.
- Automated weekly and monthly ETL/reporting workflows with SQL queries and Python transformations, reducing manual reporting effort by 25% and improving turnaround time.
- Standardized data cleaning rules, reconciled KPI definitions, and resolved missing/duplicate records across 10+ workflows, improving reporting accuracy by 15%.

PROJECTS

Telecom B2C Customer Intelligence Platform | *Python, SQL, XGBoost, SHAP, Transformers* | [Link](#)

- Built a churn and retention platform using 7,043 customers plus 142,980 usage, 84,516 billing, 3,052 incident, and 2,786 ticket records; XGBoost reached 90.46% ROC AUC and 84.46% accuracy.
- Created risk scores identifying 1,368 high-risk customers and \$229K expected monthly revenue at risk, with SHAP explanations, complaint classification, SQL views, Docker, and 21 Pytest checks.

Enterprise Urban Mobility Data Platform | *Python, PostgreSQL, dbt, Airflow, FastAPI, Power BI* | [Link](#)

- Built an end-to-end analytics platform for 10.4M+ NYC taxi records using Python ETL, PostgreSQL fact/dimension models, dbt transformations, 40 tests, and Airflow orchestration.
- Delivered FastAPI endpoints, Power BI dashboards, and ML features for demand, revenue, zone performance, payment trends, and trip-duration prediction with 0.90 R-squared.

Healthcare Time Series Forecasting for Public Health Monitoring | *Python, R, Tableau, XGBoost, LSTM* | [Link](#)

- Analyzed 20+ years of Peru public health data to identify disease trends, seasonality, regional outbreak patterns, anomalies, and high-burden populations.
- Compared ARIMA, Random Forest, XGBoost, and LSTM models with rolling validation; XGBoost reached R-squared 0.935 for key segments and Tableau dashboards supported decision-making.

Local AI Assistant with RAG and Multi-Model Inference | *Python, FastAPI, Ollama, ChromaDB, LLMs* | [Link](#)

- Built a private local RAG assistant with MiniLM embeddings, cosine search, source attribution, FastAPI endpoints, Streamlit UI, and structured JSON outputs.
- Benchmarked 3 open-source LLMs; Llama 3.2 3B reached 45 tokens/sec, 2.4x faster than Mistral 7B with 53% lower memory usage.

Amazon Sales Performance Dashboard | *Power BI, DAX, Power Query, Excel* | [Link](#)

- Created a 3-page interactive Power BI dashboard analyzing 66M sales, 1M units, 308 sellers, order status, returns, reviews, city/state performance, and product trends.
- Built DAX measures, Power Query transformations, slicers, and drill-down visuals to compare shipped, pending, cancelled, returned, and lost-order patterns across 90+ days.

EDUCATION

University of Oklahoma (GPA: 4.0/4.0)

Master of Science, Data Science and Analytics

Jawaharlal Nehru Technological University

Bachelor of Technology, Computer Science and Engineering

Aug 2024 – May 2026

Norman, OK

Dec 2020 – May 2024

Hyderabad, India